

Why AI Ethics Isn't About AI: Algorithms, Justice, and the Urban Ghetto

Alexander Williams Tolbert

1 Introduction

The increasing reliance on algorithms for decision-making in various societal sectors has raised significant concerns about their potential to reinforce existing inequalities and systemic injustices. These tools, designed to enhance efficiency and accuracy, have sparked a contentious debate due to their potential for disparate impact and reinforcement of existing inequalities (O'Neil 2016; Barocas and Selbst 2016). This paper aims to dissect the complexities surrounding the use of algorithms in societal decision-making, focusing on two case studies: law enforcement in marginalized communities and predictive policing.

Although there is a wealth of research on the negative impact of algorithms (O'Neil 2016; Barocas and Selbst 2016) and the systemic injustices experienced by marginalized groups (Wacquant 2008; Sampson 2012), little attention has been given to the intersection of these critical issues. This study aims to fill this gap by utilizing Shelby (2016)'s analysis of systemic injustices to examine the use of algorithms in societal decision-making. The central argument of this paper is that while algorithms can potentially worsen existing inequalities, their implementation can be justified as a means of protecting the most vulnerable members of society, upholding natural duties, and promoting justice. Furthermore, the challenges presented in the paradigmatic cases used to support the field of algorithmic fairness in criminal justice are ultimately rooted in political philosophy, and additional justification is needed to demonstrate what is uniquely algorithmic or what aspects of artificial intelligence make these cases particularly challenging ethically.

In order to increase clarity, this document is divided into several sections. The first section, Section 2, presents a case study on Compass and Predictive Policing. Section 3 delves deeper into the theoretical framework, examining both Rawlsian framework and Shelby's analysis of racism and discrimination. In Section 4, I argue that the paradigmatic cases used to support algorithmic fairness in criminal justice are ultimately rooted in political philosophy. Finally, Section 5 summarizes the paper's key findings.

2 Case Study: COMPAS and Predictive Policing

Algorithmic decision-making has become increasingly prevalent in the criminal justice system, using algorithms for various purposes, including risk assessment, predictive policing, and sentencing decisions. This section will focus on two widely used algorithms: COMPAS (Correctional Offender Management Profiling for Alternative Sanctions) and predictive policing algorithms.

2.1 COMPAS

COMPAS (Correctional Offender Management Profiling for Alternative Sanctions) is a widely used risk assessment tool in the United States designed to predict the likelihood of a defendant's recidivism. Developed by Northpointe Inc., now known as Equivant, COMPAS uses a proprietary algorithm that considers various factors, including criminal history, age, and gender, to assign a risk score to individuals. Judges and parole officers then use these scores to decide bail, sentencing, and parole.

However, COMPAS has been the subject of significant criticism due to its racial biases. A landmark study conducted by Angwin et al. (2016) (Angwin et al. 2016) revealed that the algorithm was almost twice as likely to incorrectly predict future criminality for black defendants than for white defendants. Conversely, white defendants were more likely to be incorrectly labeled as low-risk than black defendants. This finding raised severe concerns about the fairness and impartiality of the algorithm, as it suggested that the system was unfairly disadvantaging black defendants.

The racial bias in COMPAS's predictions is not merely a technical glitch but a manifestation of broader systemic injustices. The algorithm is trained on historical data, reflecting existing criminal justice system biases. For instance, if law enforcement has historically targeted black communities at higher rates, the data will reflect this bias, and the algorithm will learn to associate black individuals with a higher risk of recidivism. As a result, the algorithm reinforces and perpetuates these biases by making predictions that are more likely to label black defendants as high-risk. This, in turn, can lead to harsher sentences and a higher likelihood of incarceration for black defendants, exacerbating existing racial disparities in the criminal justice system.

The issues with COMPAS are not isolated but reflect broader challenges associated with using algorithms in the criminal justice system. Another area where these challenges are particularly pronounced is predictive policing, which involves using algorithms and data analytics to predict where and when crimes might occur. The following subsection will delve deeper into the issues associated with predictive policing algorithms and their potential to reinforce existing biases and exacerbate inequalities, particularly against marginalized communities.

2.2 Predictive Policing

Predictive policing involves using algorithms and data analytics to predict where and when crimes will likely occur. These algorithms use historical crime data, geographic information, and other relevant factors to identify 'hotspots' where crimes are more likely to occur. Law enforcement agencies then use these predictions to allocate resources more efficiently, target patrols, and implement preventive measures.

However, predictive policing algorithms have been criticized for reinforcing existing biases in the criminal justice system. A study by Lum and Isaac (2016) (Lum and Isaac 2016) found that predictive policing algorithms can lead to a feedback loop where police are repeatedly sent to the same neighborhoods, leading to over-policing of those areas and an increase in recorded crimes, which in turn reinforces the algorithm's predictions.

This feedback loop is particularly concerning because it can exacerbate existing inequalities. Neighborhoods already over-policed and disproportionately affected by crime will likely be further targeted by predictive policing algorithms, leading to a cycle of over-policing and increased recorded crimes. This can have a range of adverse effects on the communities targeted by these algorithms, including increased tensions between the community and the police and a higher likelihood of negative interactions with the police.

Moreover, the data used to train predictive policing algorithms often reflects existing biases in the criminal justice system. For example, suppose law enforcement has historically targeted specific neighborhoods or demographic groups at higher rates. The data will reflect this bias, and the algorithm will learn to associate those areas or groups with a higher crime risk. This can lead to a self-fulfilling prophecy, where the algorithm's predictions lead to increased police presence and surveillance in certain areas, leading to more recorded crimes and arrests, reinforcing the algorithm's predictions.

The use of predictive policing algorithms raises significant ethical and philosophical questions. For instance, is it ethical to base predictions on historical data that could be biased? Is it justifiable to increase police presence and surveillance in certain regions based on algorithm predictions, even if it results in over-policing and adverse effects on the community? These questions indicate the need for a nuanced and cautious approach to using predictive policing algorithms, considering the broader social and structural context in which they are implemented.

Furthermore, predictive policing and parole determinations based on complex algorithms make it more challenging to determine the impact of social positions on crime likelihood. While some argue that we should continue to use these methods because of the importance of knowing who is likely to engage in crime, others suggest that we should work on the social structure so that people are no longer in positions that make them more likely to engage in crime. These are difficult questions that must be addressed.

Tommie Shelby's analysis of the systemic injustices faced by ghetto residents provides valuable insight into the evaluation of predictive policing algorithms.

According to Shelby, the systemic disadvantages faced by ghetto residents fundamentally alter their obligations to the law, which challenges the prevailing discourse surrounding algorithmic fairness. This perspective emphasizes the need to consider the broader social and structural context in which these algorithms are deployed. In the next section, we will delve deeper into Shelby’s analysis and its implications for using algorithms in law enforcement, particularly in the context of the ghetto.

3 Theoretical Framework

3.1 Understanding Racism and Discrimination: Implications for Algorithmic Decision-Making

The discourse surrounding bias, especially in the context of algorithms, is multifaceted. It prompts us to question what we mean by “bias” and under which circumstances such biases become problematic. While the connection to algorithms might not seem immediate, it becomes evident when we attempt to delineate problematic biases.

Jorge Garcia’s perspective on “racism in the heart” contrasts the broader understanding of structural or systemic racism. A prevalent viewpoint posits that “racial discrimination” necessitates intentional actions by the discriminator. Such intentions could be evident in overt actions, like refusing a loan based on racial identity, or in subtler systemic actions, such as implementing business policies that disproportionately disadvantage Black individuals. However, I share a growing sentiment suggesting that policies might be labeled discriminatory based on their outcomes alone. This perspective birthed the “disparate impact” concept, recognizing the potential for structural or systemic factors to yield racialized outcomes, especially against historical racial biases.

Debates over the terminology in discussions about racism and discrimination are rife. Some advocate using terms like “racist” or “discriminatory” exclusively for intentional actions, fearing that a broader application might insinuate moral blame or condemnation. However, a pressing normative question emerges: Should we confront structural outcomes that arise without any individual’s racist intent? This dilemma resonates in legal circles, where the “disparate impact” doctrine has evolved to adopt a non-moralized stance. A prevalent ‘no fault’ liability model has surfaced, allowing entities to acknowledge the unintended discriminatory effects of their policies and express a willingness to amend them. If they view such policies as essential, they might advocate for structural modifications at different stages of the process.

Understanding the nuances of racism and discrimination becomes crucial when evaluating the ethical challenges posed by algorithmic decision-making. Without careful design, algorithms risk magnifying existing biases and disparities. This section delves into the insights of Jorge Garcia and Tommie Shelby, exploring their implications for ethically deploying algorithms in societal contexts.

Jorge Garcia elucidates the noncognitive dimensions of racism, categorizing these attitudes into three primary groups: states of desire, volitional states, and affective states. States of desire encompass wishes for a specific racial group to face adversity or indifference to their plight. Volitional states relate to decisions or intentions to harm a racial group or a lack of effort to avert such harm. Affective states span emotions ranging from contempt to callous indifference directed at a specific racial group. Garcia coins the term 'racial disregard' to encapsulate these attitudes, which include race-based ill will, insufficient goodwill, and disrespect, provided they are racialized (Garcia 1996).

In contrast, Tommie Shelby offers a perspective grounded in institutional practices. He argues that racial discrimination manifests when a racial characteristic, or a set of characteristics attributed to a social group, is wrongly treated as a sign of disvalue, incompetence, or inferiority (Shelby 2007, 131). This viewpoint aligns with systemic or structural racism, emphasizing how societal institutions maintain racial hierarchies and inequalities. Drawing from John Rawls's concept of 'formal justice,' Shelby posits that *institutional racism* arises when the rules of significant societal institutions consistently reflect the racial biases of those in power (131).

Shelby's discourse on racial discrimination anchors in a Rawlsian framework, emphasizing fairness, equality, and individual rights. Rawls envisions a just society where liberties, opportunities, and resources are distributed to ensure everyone's fair chance at success. Building on this, Shelby argues that systemic injustices, especially those confronting marginalized groups, challenge their civic duties and necessitate a reevaluation of justice principles in these contexts (131–132).

Furthermore, Shelby underscores the historical context of racial discrimination. Despite the abolition of overt white supremacist principles, race-based slavery, and de jure racial segregation, covert racial prejudice persists, undermining the impartial administration of major societal institutions. He contends that ongoing racial discrimination in employment and housing disproportionately affects the less affluent members of racial minority groups. The education, financial resources, and social capital of more privileged members enable them to achieve greater social mobility despite the persistence of racism (132).

In conclusion, as we delve deeper into the Rawlsian foundation of Shelby's arguments in subsequent sections, it becomes evident that comprehending the intricacies of racism and discrimination is pivotal for addressing the ethical challenges posed by algorithmic interventions in societal decision-making.

3.2 Rawlsian Framework and Shelby's Analysis

John Rawls's theory of justice as fairness provides a foundational framework for understanding and addressing social justice issues. Rawls argues that a just society is one in which the distribution of liberties, opportunities, and material advantages is organized to give everyone a fair chance to flourish. This entails ensuring that no citizen's life prospects are diminished because the social scheme disadvantages them in ways that cannot be justified on impartial grounds (Rawls

1971).

In his work “Justice, Deviance, and the Dark Ghetto”, (Shelby 2007, pp. 126-160) Tommie Shelby builds on this Rawlsian framework to provide a thought-provoking analysis of the systemic injustices that ghetto residents experience and how these injustices impact their civic duties. Shelby contends that the systemic injustices these communities face compromise their civic obligations and necessitate a reevaluation of the principles of justice in these contexts (131–132). He argues that the residents of urban ghettos, who are often subjected to ongoing racial discrimination, economic marginalization, and state-sanctioned violence, have a legitimate claim to a different set of civic obligations than those that apply to the broader society.

He notes that the residents of urban ghettos often face a unique set of challenges that are not adequately addressed by the traditional Rawlsian framework, which assumes a relatively stable and well-ordered society. Shelby argues that in contexts where systemic injustices are deeply entrenched, a more radical approach may be necessary to address these injustices’ root causes and promote a more just and equitable society (132).

Shelby elaborates on ‘deviance as a response to the ghetto plight’ (134). He defines *ghettos* by three core characteristics: race, urban location, and poverty, and argues that the issues of justice raised by U.S. ghettos are especially acute and, in some ways, unique. Shelby emphasizes that while many who live under ghetto conditions respect the law and accept conventional morality, many are alienated, some profoundly so. This alienation, combined with the desperation and shame created by ghetto poverty, leads some residents to resort to crime. Shelby contends that this criminal activity might express something more, or something other, than a character flaw or a disregard for the authority of morality (135–136).

Shelby further explores the racial significance of deviance in ghetto communities, arguing that racial prejudice continues to have a negative impact on the life chances of racial minorities in the United States, especially black citizens. He notes that the impact of institutional racism is deepest in dark ghettos, where racism and extreme poverty combine to create a uniquely stigmatized subgroup of the black population. Shelby argues that this dynamic, combined with the ghetto subculture, plays out in many arenas, including employment, housing, and the criminal justice system, and greatly disadvantages those from the ghetto who seek employment, decent housing, and good schools (139–140). Shelby also discusses the impact of street crime norms on ghetto residents who want to avoid criminal deviance. He notes that the widespread use of guns among drug dealers and muggers creates a demand for these weapons in ghetto neighborhoods, leading many residents, including children, to arm themselves for protection. This creates a sense of danger and distrust in ghetto neighborhoods, making it difficult for a broad sense of community to form or be maintained (138).

James Forman Jr.’s “Locking Up Our Own” (Forman 2017) provides a crucial historical and normative perspective that complements Shelby’s analysis. In *Locking Up Our Own: Crime and Punishment in Black America*, James Forman Jr. provides a detailed and empathetic analysis of the difficult decisions faced

by Black leaders during the 1970s and 1980s. Forman does not justify the policies that contributed to mass incarceration but provides essential context for understanding the decisions made by Black leaders at the time. Forman begins by detailing the demands from Black community members for their leaders to address the rising crime and drug addiction in their neighborhoods. He describes a scene in which a Black mother testifies before the city council in Washington D.C., pleading for them to do something about the open-air drug markets near her home (Forman 2017, p. 45).

This example illustrates the deep concern many Black residents, particularly older ones, had about the impact of crime and drugs on their communities. They supported tough-on-crime policies as a way to regain control of their neighborhoods. This demand for action must be understood in the broader historical context. Many Black leaders during this period were part of the civil rights generation and had fought for the desegregation of schools, public accommodations, and the criminal justice system. These leaders, having witnessed and experienced the racism of the criminal justice system firsthand, were often profoundly skeptical of the police and the courts. However, they felt responsible for protecting their communities from the genuine threats of crime and drug addiction. This created a tension between the desire to address these issues and the skepticism of the criminal justice system (p. 89). This creates a double bind. The political climate of the time further complicated matters. Law and order was a dominant theme, and many Black leaders felt pressure to conform to this prevailing narrative to gain political legitimacy and to show that they were as tough on crime as their white counterparts. Forman describes how Eric Holder, a young prosecutor in Washington D.C., supported mandatory minimum sentences for drug offenses, a policy he later regretted and criticized as Attorney General (p. 134). Forman's analysis provides a nuanced and empathetic understanding of the difficult decisions faced by Black leaders during this period. While he does not justify the policies that contributed to mass incarceration, he does provide important context for understanding the decisions made by Black leaders at the time.

Shelby also delves into whether the deviant conduct and attitudes in the ghetto are unreasonable. He distinguishes between three possible assessments of the basic structure of U.S. society: fully just, not fundamentally unjust, but with some injustices that should be addressed, and fundamentally unjust, requiring radical reform. He argues that the obligations of the ghetto poor would differ depending on which assessment is correct. Shelby invokes the distinction between civic obligations and natural duties, arguing that the former is owed to those with whom one is cooperating to maintain a fair basic structure. At the same time, the latter are unconditionally binding on all persons. He contends that the constitutional essentials, such as fundamental rights and political procedures, do not include fair equality of opportunity or the difference principle and that the standard for tolerable injustice should include an adequate opportunity to develop marketable skills. Shelby argues that the existence of the dark ghetto is incompatible with any meaningful form of reciprocity among free and equal citizens and that the ghetto poor are not violating the principle of

reciprocity or shirking valid civic obligations when they refuse to accept menial jobs or respect the authority of the law (Shelby 2007, 143–150).

He argues even in a fundamentally unjust society, the ghetto poor still have moral duties to one another and others. These duties, which include not being cruel, helping the needy and vulnerable, not causing unnecessary suffering, and showing mutual respect, are not negated by the existence of an unjust social order (Shelby 2007). Shelby contends that these duties make some deviant attitudes and actions of the ghetto poor impermissible, not because they are forbidden by law but because they cannot be fully justified from a moral point of view.

He argues that this duty requires individuals to support and comply with just institutions and, where just institutions do not exist, to help bring them about. This duty applies to everyone, including the oppressed and the ghetto poor. It is not based on civic reciprocity but on each person's status as a moral person (Shelby 2007). Shelby also discusses the different ways a society might be unjust and the appropriate forms of resistance in each case. He argues that the form of resistance should consider the reasonableness of existing aspects of the social scheme and choose measures of opposition accordingly. Shelby emphasizes the importance of including the ghetto poor in the resistance effort and the need to build bonds of political solidarity with each other and progressive allies. He argues that an outlaw subculture in the ghetto erodes the chance of developing mutual trust, empathy, compassion, and broader forms of loyalty, all necessary for political solidarity and collective action (Shelby 2007).

Shelby concludes by emphasizing the importance of building bonds of political solidarity among the ghetto poor and with progressive allies. He argues that an outlaw subculture cannot reign in the ghetto if such solidarity is to form and be sustained. He suggests that while the ghetto poor may be justified in engaging in street crime as a survival tactic or expression of justified rebellion, they should not regard gangsterism and hustling as vocations or allow these practices to corrupt their character (159). Instead, they should view them as survival tactics, means of self-defense, or expressions of justified rebellion, and whenever possible, publicly register their principled dissatisfaction with the existing social order (Shelby 2007). Shelby argues that even in a fundamentally unjust society, individuals, including the ghetto poor, have natural duties not negated by an unjust social order. These duties include not being cruel, helping the needy and vulnerable when it is not too risky, not causing unnecessary suffering, and showing mutual respect. These duties, Shelby emphasizes, are not suspended or void because one is oppressed (151). Shelby also discusses the duty of justice, as outlined by John Rawls, which requires individuals to support and comply with just institutions and, where just institutions do not exist, to help bring them about. He argues that this duty applies to everyone, including the oppressed and the ghetto poor, and is not based on the principle of civic reciprocity but is a duty each has as a moral person (152). He distinguishes between two ways a society might be unjust: when it fails to live up to its professed ideals or when the dominant conception of justice is an ideology that legitimates and sustains an oppressive regime. He argues that the form of

opposition should consider the reasonableness of existing aspects of the social scheme and the need for public order (Shelby 2007, 157).

Shelby's analysis, enriched by Forman's historical and normative context, presents a detailed examination of the ethical considerations surrounding using algorithms for law enforcement purposes, especially within the ghettos. Their combined perspective provides a robust foundation upon which the justifiability of these algorithms can be assessed, factoring in crimes, the broader societal backdrop, and implications for the communities in question. At the heart of this discourse lies the very definition of civic responsibilities. Shelby asserts that civic obligations are inherently based on reciprocity with the legal system. Given the systemic injustices faced by ghetto residents, it stands to reason that this reciprocal relationship becomes fundamentally skewed. If the state fails in its duties to specific citizens, is it not conceivable that these citizens' obligations toward the state are, in turn, diminished? Shelby delves deeper, pointing out that while systemic inequalities might excuse some civic obligations, they do not negate natural duties, which all individuals bear regardless of their circumstances. Essential to this discourse are serious offenses such as assault, rape, or murder. Shelby argues that even within contexts of deep-seated injustice, individuals retain a natural duty to refrain from these grave offenses.

In light of this, predictive policing algorithms, when applied within ghettos, serve a dual purpose. Firstly, they ensure individuals adhere to these natural duties by targeting violent crimes. Secondly, they can serve a broader societal objective: protecting those already systematically marginalized. Using Shelby's Rawlsian lens of fairness, equality, and individual rights, it becomes evident that the primary obligation of a just legal system is to shield its most vulnerable. In the context of ghettos, this duty is paramount. The ghettos, shaped by structural inequalities of American institutions, become areas where the residents, already grappling with systemic disadvantages, need heightened protection. Violent crimes inflict long-lasting damage beyond immediate harm by creating a pervasive environment of fear and vulnerability. Thus, if a predictive algorithm is adept at pinpointing potential violent offenders within the ghetto, it arguably aligns with the principles of justice, safeguarding the community's most vulnerable.

However, this justification becomes convoluted when the algorithm's focus shifts. If it predominantly targets crimes committed against individuals "outside the ghetto," it risks perpetuating a narrative that places the safety and interests of the more privileged above those of ghetto residents. "Outside the ghetto" denotes regions not marred by concentrated poverty, systemic discrimination, and racial segregation, which define ghettos. Inhabitants of these areas, often privileged regarding race, class, and socio-economic status, navigate different social dynamics.

Using algorithms that seemingly prioritize these external regions over the ghettos is problematic. Not only does it reinforce power imbalances, but it also subtly insinuates that the systemic challenges within ghettos are secondary. By focusing on crimes against those already advantaged, the justice system might inadvertently cement the existing hierarchies rather than striving to dismantle

them. Further, suppose an algorithm targets individuals committing non-violent offenses or minor infractions, such as violating open container laws or smoking cannabis. In that case, this raises significant philosophical questions about the nature of crime, the purpose of bail, and the criteria that should be employed to make bail decisions. If the algorithm effectively fulfills its intended purpose of preventing violent crime or ensuring public safety, then its utilization could be justified. The pivotal philosophical question, in this case, centers on the types of arrests that should be classified as violent. This is predominantly a political philosophy issue rather than an issue intrinsic to the algorithm.

4 Reduction to Political Philosophy

The examples of COMPAS and predictive policing are frequently used to highlight the potential for AI to magnify and perpetuate biases, sparking conversations about fairness in AI. However, a closer examination reveals these issues are more than just algorithmic. Instead, they are deeply intertwined with broader political and philosophical questions surrounding justice, societal responsibility, and the very nature of crime. As Shelby and Forman Jr. have noted, while predictive policing tools may be accurate or biased, the underlying dilemmas are rooted in longstanding societal structures and systemic injustices that must be addressed to create a truly equitable and just criminal justice system.

The COMPAS and predictive policing cases often show how AI can perpetuate and amplify existing biases, making them paradigmatic in discussions about AI fairness. However, it becomes clear that the core issues at play are not inherently tied to AI. Instead, they are deeply rooted in longstanding societal structures and systemic injustices. In criminal justice, there is a central dilemma between ensuring immediate public safety and addressing the root causes of crime, which can be called corrective justice. This trade-off is present regardless of the presence or perfection of predictive algorithms. The real question is how society can balance the urgency of immediate safety against the long-term need to address the root causes of crime.

Both Shelby and Forman Jr. emphasize that the nature, justification, and societal response to crime can vary based on its context (Shelby 2007; Forman 2017). Factors such as the perpetrator, the location of the crime, and the victim play a pivotal role in shaping this context. This understanding underscores that the challenges presented by crime are not merely technical but are profoundly political and philosophical. Given individuals' systemic injustices, some actions labeled as "crimes" might be viewed differently in specific contexts. It is worth noting that crimes in marginalized areas may be treated differently than those in affluent areas, even if the nature of the crime is identical. Shelby's analysis suggests that these actions might be justified or understood differently, challenging the foundational societal views on crime and justice (Shelby 2007). Determining what constitutes bias or unjust discrimination in the context of AI is a central challenge. Not all biases are inherently detrimental; some may be grounded in factual data. However, when these biases perpetuate systemic injustices or are

based on unfounded stereotypes, they become problematic. The use of AI in predictive policing accentuates these questions: When does an algorithm’s prediction, based on historical data, translate into unjust discrimination? These are ultimately political philosophy issues for philosophers of race and political philosophy more broadly, irrespective of the facts of AI.

Historical insights show that these challenges predate AI, especially within Black communities. The dilemmas faced by Black communities, especially the trade-offs between public safety and justice, have historical roots. Forman highlights that these challenges existed and were addressed by Black leaders without the influence of algorithms. They were grappling with real-world pressures, societal structures, and the immediate threats of their time. This context emphasizes that the ethical dilemmas we associate with AI today are not new but are continuations of longstanding societal debates. The decisions made by Black leaders in the past were influenced by the tangible challenges of their era, not by algorithms. Further insights suggest that the justification for using predictive algorithms is not about the technology but the broader societal context in which it operates. The core of my argument is that the dilemmas associated with predictive policing exist independently of the algorithms. While tools like COMPAS bring these dilemmas to the forefront, they are not uniquely AI issues. The real challenges lie in broader societal structures, values, and decisions.

Bias Amplification is a significant issue in AI. It can amplify existing biases, but the root cause of these biases lies in historical and systemic discrimination. The algorithms are trained on data reflecting societal prejudices, perpetuating what they are fed. Feedback Loops are also worrying. Predictive policing algorithms create feedback loops that reinforce crime predictions. However, the over-policing of specific neighborhoods and the subsequent reinforcement of crime predictions is a manifestation of longstanding discriminatory practices, not a novel challenge introduced by AI. Opacity and Accountability are often cited as distinctive challenges of AI models. However, the real issue is not the algorithm’s opacity but the lack of Accountability in the systems that deploy them. This lack of Accountability has been a feature of many institutional practices, even before the advent of AI.

Ensuring the safety of a community while preserving individual rights presents a significant challenge. Predictive algorithms can prioritize public safety by forecasting potential future crimes, but relying on broad statistical data instead of individualized evidence may infringe on individual rights. While it is essential to address immediate threats to safety, it is equally important to address the systemic issues leading to those threats. Addressing the root causes of crime may provide a more sustainable solution than merely predicting and punishing potential future criminals. There is an inherent tension between the desire to address and correct past injustices (corrective or reparative justice) and the need to ensure public safety. Predictive algorithms lean towards public safety by trying to forecast future crimes. However, they do not account for the broader societal structures that place specific individuals at a higher risk of committing crimes, not due to their moral character but their circumstances. In that case, they can inadvertently perpetuate systemic issues.

We face a conundrum where, on the one hand, there is a genuine need to identify potential future criminals for the sake of public safety. On the other hand, we recognize that specific individuals are statistically more likely to commit crimes due to systemic issues, not personal failings. The challenge is using these predictive tools without exacerbating or perpetuating these systemic problems. The debate around algorithmic bias is central to this discussion. If an algorithm reflects societal biases, is it biased or just a mirror of our prejudices? The real challenge is not just about "fixing" the algorithm but addressing the root societal issues that the algorithm brings to the forefront. My argument underscores that while AI and algorithms like COMPAS are modern tools, the ethical and societal dilemmas they present are not unique to the AI era. Instead, they are the latest manifestation of longstanding political and philosophical debates about justice, public safety, and societal structures. The real challenge in AI ethics is distinguishing what is genuinely a new ethical dilemma introduced by AI and what is a continuation of older, deeper societal issues.

The debate around AI ethics and algorithmic fairness revolves around identifying unique ethical challenges that arise in the context of AI. Although these ethical dilemmas existed long before AI and would persist even in its absence, it is crucial to comprehend how AI reflects and amplifies societal structures and values. This examination requires a deep dive into the heart of ethical debates, asking whether the technology or societal structures pose the primary challenge. Using AI tools, like predictive policing algorithms such as COMPAS, raises ethical dilemmas reflecting longstanding political philosophy debates. These debates involve balancing public safety and individual rights, immediate threats and systemic issues, societal values, and systemic inequalities. They also encompass identity issues, such as the identity of the perpetrator and the victim. It is important to note that these challenges existed before AI, especially within Black communities. The decisions made by Black leaders in the past were influenced by the tangible challenges of their time, not algorithms. While predictive algorithms may be justifiable in specific contexts, the broader societal context in which they operate is critical.

The location and nature of the crime, as well as the identity dynamics of the perpetrator and victim, are critical considerations. Questions about where crimes occur (inside or outside the ghetto) and what crimes are being committed are essential. Racial, socioeconomic, and community dynamics are factors to consider. An algorithm in specific contexts, like the ghetto, can be justified based on various factors. However, these justifications are not about the algorithm itself. Whether the algorithm is present, functions flawlessly, or is absent, these foundational questions persist. The discourse on AI ethics and algorithmic fairness faces the challenge of identifying the unique ethical dilemmas posed by AI. What novel issues does AI bring if these dilemmas existed before AI and would continue without it? Even with a "perfect" predictive policing algorithm that could predict crime with complete accuracy, ethical questions would persist. Who gets targeted, what crimes are prioritized, and how do we balance public safety with individual rights? These questions are political and philosophical, not technical.

5 Conclusion

This paper has delved into the complexities surrounding using algorithms in law enforcement within the ghetto, examining the associated challenges, ethical dilemmas, and potential justifications. We have dissected the potential of algorithms to both exacerbate existing inequalities and reinforce power dynamics and their capacity to safeguard vulnerable community members, uphold natural duties, and advance justice. This intricate interplay of conflicting demands between ethical or justice-related considerations raises a pivotal question: Are there scenarios within the realm of algorithms that necessitate making trade-offs between conflicting ethical and justice-related demands?

Many trade-offs encountered in algorithmic fairness can be traced back to more general dilemmas in ethics or political philosophy. This observation prompts a critical inquiry into whether any trade-offs are unique to algorithmic fairness or whether all such dilemmas ultimately reduce to more general ethical or political challenges. Furthermore, it encourages us to reflect on what constitutes genuine instances of trade-offs, even if they are not exclusive to algorithmic fairness.

A particularly challenging ethical trade-off arises with deep learning algorithms, which often operate invisibly without our knowledge and are inherently difficult to comprehend. Unlike other technologies that require active choice or consent, and thus awareness of their presence and use, algorithms and machine learning often operate invisibly, raising ethical dilemmas related to consent, autonomy, and transparency.

In conclusion, deploying algorithms in law enforcement within the ghetto presents many ethical and justice-related dilemmas. While deploying algorithms can be justified to safeguard the most vulnerable community members, uphold natural duties, and advance justice, it also harbors the potential to exacerbate existing inequalities and reinforce prevailing power dynamics. It is imperative to actively involve affected communities in the decision-making process, ensure greater transparency and accountability, address broader social and structural injustices, and continuously tie back the discussion to the central contention of this paper. Only through a meticulous and comprehensive consideration of these challenges can we hope to promote justice and fairness in using algorithms in law enforcement within the ghetto.

References

- Angwin, Julia, Jeff Larson, Surya Mattu, and Lauren Kirchner. 2016. “Machine bias: There’s software used across the country to predict future criminals. And it’s biased against blacks.” *ProPublica*.
- Barocas, Solon, and Andrew D. Selbst. 2016. “Big Data’s Disparate Impact.” *California Law Review* 104:671–732.

- Forman, James. 2017. *Locking Up Our Own: Crime and Punishment in Black America*. Farrar, Straus / Giroux. ISBN: 978-0374189976.
- Garcia, Jorge L. A. 1996. "The Heart of Racism." *Journal of Social Philosophy* 27 (1): 5–45.
- Lum, Kristian, and William Isaac. 2016. "To predict and serve?" *Significance* 13 (5): 14–19.
- O'Neil, Cathy. 2016. *Weapons of Math Destruction: How Big Data Increases Inequality and Threatens Democracy*. Crown.
- Rawls, John. 1971. *A Theory of Justice*. Harvard University Press.
- Sampson, Robert J. 2012. "Great American City: Chicago and the Enduring Neighborhood Effect."
- Shelby, Tommie. 2007. "Justice, deviance, and the dark ghetto." *Philosophy & Public Affairs* 35 (2): 126–160.
- . 2016. *Dark Ghettos: Injustice, Dissent, and Reform*. Harvard University Press.
- Wacquant, Loïc. 2008. *Urban Outcasts: A Comparative Sociology of Advanced Marginality*. Polity.